GERMÁN KÜBER

Senior Backend Engineer — Production Python — Real-Time AI & Voice Pipelines

Las Palmas, Spain (EU Citizen) — Remote — EU / Americas (UTC-5 to UTC+2) — +34 666 844 731
german.kuber@gmail.com — linkedin.com/in/germankuber — github.com/GermanKuber — germankuber.com

PROFESSIONAL SUMMARY

Senior backend engineer with 15+ years building production systems and a deep focus over the last several years on production Python (FastAPI, Django, Flask) for real-time AI products. Shipped a real-time audio AI pipeline at KliverAI on OpenAI Realtime API with sub-200ms streaming latency and a WebSocket-based transcription pipeline capturing live meeting audio — directly applicable to ambient AI / scribe products. Operated production Python and Rust services on Docker and Kubernetes / AWS at high scale (Borderless Capital: $50M+ daily transactional flow). Comfortable across adjacent stacks including **TypeScript / Node.js** and .NET. Customer-facing impact orientation; comfortable with ambiguity and fast iteration; very much a fit for ship-it-now environments.

CORE EXPERTISE

Production Python: FastAPI, Flask, Django, asyncio, async / await, Cython, PyO3 bindings, Pydantic, SQLAlchemy, REST and WebSocket APIs, background workers, structured logging, testing (pytest)

Real-Time AI & Voice Pipelines: OpenAI Realtime API, WebSocket streaming, SignalR, sub-200ms audio pipelines, real-time transcription capture, real-time evaluation pipelines, LLM orchestration (OpenAI, Anthropic, Gemini), multi-provider routing, prompt engineering, RAG, multi-agent (LangGraph), MCP server / client

Cloud & Infrastructure: AWS, Docker, Kubernetes (manifests, deployments, services), CI / CD pipelines, Linux, Infrastructure as Code (Terraform), observability stacks

Data & Storage: PostgreSQL (OLTP, jsonb, semantic cache backing store), caching and pub/sub (Redis-style), high-throughput data pipelines

Backend Engineering: Event-driven microservices, async coordination, fault tolerance, robust APIs and abstractions, Domain-Driven Design, customer-facing systems at scale

Observability & Reliability: Prometheus, OpenTelemetry, distributed tracing, structured logging, 50+ production metrics, 24 / 7 on-call rotation, post-mortems, incident response

Adjacent Stacks: TypeScript / Node.js, Rust (tokio, async), .NET / C#, Go (reverse-engineering Geth internals)

PROFESSIONAL EXPERIENCE

Lead AI Architect & Co-Founder

KliverAI (Own Venture, closed Apr 2026) — Apr 2025 – Apr 2026 (full-time) — Remote, Spain

- Designed and shipped a real-time audio AI pipeline on OpenAI Realtime API with SignalR streaming, achieving sub-200ms end-to-end latency across concurrent user sessions — the same class of problem as an ambient AI scribe: capturing live conversation, processing it through LLMs, and returning structured output in real time.
- Built a real-time transcription capture pipeline: a Chrome Extension streams live Google Meet audio over WebSocket into a Python backend, which then feeds an automated AI evaluation pipeline.
- Built backend services in Python (FastAPI, asyncio) for AI integrations and data pipelines, alongside Node.js / TypeScript services for the API and streaming layer; deployed everything on Docker and Kubernetes on AWS.
- Integrated three LLM providers (OpenAI, Anthropic, Gemini) with dynamic model routing for cost and latency optimization.
- Built an adaptive multi-agent platform on LangGraph with multi-step branching logic and RAG-powered personalization.
- Implemented MCP (Model Context Protocol) servers and clients to standardize tool integration across services.
- Established the observability stack with OpenTelemetry, distributed tracing, structured logging, and LLM-specific tracing (Langfuse, LangSmith) plus prompt evaluation (Braintrust, PromptLayer, Promptfoo).
- Hired and led a distributed team of up to 3 engineers, owning technical interviews, hiring, and architectural reviews — hands-on in the codebase throughout.

Fractional CTO (Part-time, hands-on)

Rocking Product — Apr 2025 – Nov 2025 (part-time, parallel to KliverAI) — Remote / Buenos Aires

- Shipped customer-facing AI features in production using Python (FastAPI, LangChain) and Node.js / TypeScript, establishing the team's first LLM evaluation and monitoring practices.
- Optimized LLM API usage and cloud infrastructure (AWS, Docker, Kubernetes) through model selection, prompt compression, and semantic caching — cutting combined LLM API and cloud spend by 45%.
- Inherited an 8-person engineering team across backend, frontend, and DevOps; stabilized release cadence and reduced sprint carryover within the first quarter.
- Established code review standards, CI / CD pipelines, and engineering rituals while remaining hands-on in the codebase.
- Partnered directly with founders and clients to translate business priorities into technical roadmaps — typical early-stage, ambiguous, ship-it-now environment.

Lead Backend Engineer (Distributed Systems, Python & Rust)
Borderless Capital / CTF Capital — Aug 2021 – Apr 2025 — Remote

- Operated production backend services processing $50M+ in daily transactional flow and contributing $100K+ in attributable annual P&L with execution under 10ms p99 latency — real users, real money, real consequences if anything broke.
- Built Python services (endpoints, business logic, external integrations) running on Docker and Kubernetes in production, alongside the Rust core; owned K8s manifests and deployment configuration for the services I shipped.
- Built Python / Cython and PyO3 FFI bindings exposing Rust performance-critical code to higher-level Python orchestration layers, eliminating GC pauses across language boundaries.
- Used Postgres as the system of record for pipeline state, opportunities, and execution outcomes, and Redis-style caching both for low-latency caching and as a pub / sub backbone between services.
- Engineered low-latency observability with Prometheus and custom telemetry, tracking 50+ real-time metrics across the production pipeline.
- Owned 24 / 7 production reliability and incident response, including on-call rotation, post-mortems, and observability ownership.
- Acted as technical lead: architecture, code review, mentoring, and hiring for backend engineers across Python and Rust.

Senior Software Architect (Backend, Event-Driven)
Paramo Technology — Aug 2020 – Aug 2021 — Remote / Argentina

- Held company-wide architecture authority for a large-scale online casino platform — accountable for cross-team technical decisions across 9+ engineering teams (200+ professionals, each with its own coordinator), working directly with the CTO and defining organization-wide engineering standards.
- Owned end-to-end backend architecture with event-driven microservices, real-time state management, and strict consistency requirements.
- Led technical interviewing and hiring across the engineering organization; delivered internal training on architecture and operational practices.

Software Architect (Full-Stack)
CloudX (Software Consultancy) — Aug 2017 – Aug 2020 — Buenos Aires

- Led architecture decisions across multiple enterprise product lines, designing microservices backends with Domain-Driven Design.
- Delivered technical training and conference talks on backend architecture and software design.

Software Architect
Majestic Hotel Group — Jan 2014 – Aug 2017 — Argentina

- Designed and shipped full-stack features for internal hotel management platforms; refactored backend toward an event-oriented architecture.

.NET Web Developer
Plus54 — Aug 2009 – Mar 2014 — Buenos Aires

- Early career: built .NET web applications, internal tooling, and mobile components end to end.

KEY ACHIEVEMENTS

Real-time audio AI pipeline at KliverAI — OpenAI Realtime API + SignalR + WebSocket transcription capture, sub-200ms streaming latency across concurrent sessions.
Production backend at Borderless — Python and Rust services on Docker / Kubernetes / AWS, processing $50M+ in daily transactional flow and contributing $100K+ in attributable annual P&L at sub-10ms p99.
Microsoft MVP, Developer Technologies — 8 consecutive years (2017–present).
80+ International Technical Talks — conferences and meetups across LATAM and Europe.
Net-Baires Founder — Argentina's largest .NET developer community, 100K+ members across 10+ years.

EDUCATION & CERTIFICATIONS

Microsoft Certified: Azure Solutions Architect Expert (AZ-305) — Microsoft
Microsoft Certified: Azure Associate (AZ-104 / AZ-204 level) — Microsoft
Microsoft MVP, Developer Technologies — Feb 2017 – Present (8 consecutive years)
Systems Engineer — Universidad Abierta Interamericana (UAI), Buenos Aires, Argentina (2017–2021)
Polkadot Builder Program — Polkadot Blockchain Academy, Lisbon, Portugal (2026)

OPEN SOURCE, WRITING & COMMUNITY

YouTube channel (youtube.com/@germankuber) — technical deep-dives on AI, Python, Rust, and software architecture.
Net-Baires (.NET Community) — Founder (2015–present). 100K+ members.
Crypto Para Mortales — Founder (2023–present). Spanish-language blockchain and crypto education conference.
Patrones Para Mortales — Founder (2020–present). Educational content on software design.

SPOKEN LANGUAGES

Spanish: Native — English: Full Professional Proficiency — daily use in international remote teams